

---

# GPU programming: Code optimization part 3 and advanced features

---

Sylvain Collange  
Inria Rennes – Bretagne Atlantique  
[sylvain.collange@inria.fr](mailto:sylvain.collange@inria.fr)

# Outline

---

- Instruction-level optimization
  - ◆ A few device code features
  - ◆ Warp divergence
- Advanced features
  - ◆ Device-mapped memory, Unified virtual address space, Unified memory
  - ◆ Global and shared memory atomics
  - ◆ Warp-synchronous programming, warp vote, shuffle

# Device functions

- Kernel can call functions
- Need to be marked for GPU compilation

```
__device__ int foo(int i) {  
}
```

- A function can be compiled for both host and device

```
__host__ __device__ int bar(int i) {  
}
```

- Device functions can call device functions
  - ◆ Older GPUs do not support recursion

# Local memory

- Registers are fast but
  - ◆ Limited in size
  - ◆ Not addressable
- Local memory used for
  - ◆ Local variables that do not fit in registers (*register spilling*)
  - ◆ Local arrays accessed with indirection

```
int a[17];  
b = a[i];
```

- **Warning:** local is a misnomer!
  - ◆ Physically, local memory usually goes off-chip
  - ◆ About same performance as coalesced access to global memory

# Loop unrolling

- Can improve performance
  - ◆ Amortizes loop overhead over several iterations
  - ◆ May allow constant propagation, common sub-expression elimination...
- Unrolling is **necessary** to keep arrays in registers

Not unrolled

```
int a[4];
for(int i = 0; i < 4; i++) {
    a[i] = 3 * i;
}
```

Indirect addressing:  
a in local memory

Unrolled

```
int a[4];
a[0] = 3 * 0;
a[1] = 3 * 1;
a[2] = 3 * 2;
a[3] = 3 * 3;
```

Static addressing:  
a in registers

Trivial  
computations:  
optimized away

- The compiler can unroll for you

```
#pragma unroll
for(int i = 0; i < 4; i++) {
    a[i] = 3 * i;
}
```

# Device-side functions: C library support

- Extensive math function support
  - ◆ Standard C99 math functions in single and double precision:  
e.g. `sqrtf`, `sqrt`, `expf`, `exp`, `sinf`, `sin`, `erf`, `j0`...
  - ◆ More exotic math functions:  
`cospi`, `erfcinv`, `normcdf`...
  - ◆ Complete list with error bounds in the CUDA C programming guide
- Starting from CC 2.0
  - ◆ `printf`
  - ◆ `memcpy`, `memset`
  - ◆ `malloc`, `free`: allocate global memory on demand

# Device-side intrinsics

- C functions that translate to one/a few machine instructions
  - ◆ Access features not easily expressed in C
- Hardware math functions in single-precision
  - ◆ GPUs have dedicated hardware for reverse square root, inverse, log2, exp2, sin, cos in single precision
  - ◆ Intrinsics: `__rsqrt`, `__rcp`, `exp2f`, `__expf`, `__exp10f`, `__log2f`, `__logf`, `__log10f`, `__sinf`, `__cosf`, `__sincosf`, `__tanf`, `__powf`
  - ◆ Less accurate than software implementation (except `exp2f`), but much faster
- Optimized arithmetic functions
  - ◆ `__brev`, `__popc`, `__clz`, `__ffs`...  
Bit reversal, population count, count leading zeroes, find first bit set...
  - ◆ Check the CUDA Toolkit Reference Manual
- We will see other intrinsics in the next part

# Outline

---

- Instruction-level optimization
  - ◆ A few device code features
  - ◆ Warp divergence
- Advanced features
  - ◆ Device-mapped memory, Unified virtual address space, Unified memory
  - ◆ Global and shared memory atomics
  - ◆ Warp-synchronous programming, warp vote, shuffle



# Warp-based execution

---

## Reminder



- Threads in a warp run in lockstep
- On current NVIDIA architectures, warp is 32 threads
- A block is made of warps
  - ◆ Warps do not cross block boundaries
  - ◆ Block size multiple of 32 for best performance

# Branch divergence

- Conditional block

```
if(c) {  
    // A  
}  
else {  
    // B  
}
```

- When all threads of a warp take the same path:




With imaginary 4-thread warps

# Branch divergence

- Conditional block

```
if(c) {  
    // A  
}  
else {  
    // B  
}
```

- When threads in a warp take different paths:



- Warps have to go through both **A** and **B**: lower performance

# Avoiding branch divergence

- Hoist identical computations and memory accesses outside conditional blocks

```
if(tid % 2) {  
    s += 1.0f/tid;  
}  
else {  
    s -= 1.0f/tid;  
}
```

```
float t = 1.0f/tid;  
if(tid % 2) {  
    s += t;  
}  
else {  
    s -= t;  
}
```


- When possible, re-schedule work to make non-divergent warps

```
// Compute 2 values per thread  
int i = 2 * tid;  
s += 1.0f/i - 1.0f/(i+1);
```

- What if I use C's ternary operator (?:) instead of if?  
(or tricks like ANDing with a mask, multiplying by a boolean...)

# Answer: ternary operator or predication

- Run both branches and select:  $R = c ? A : B;$ 
  - ◆ No more divergence?
- All threads have to take both paths  
No matter whether the condition is divergent or not




- Does **not** solve divergence: we lose in all cases!
- Only benefit: fewer instructions
  - ◆ May be faster for short, often-divergent branches
- Compiler will choose automatically when to use predication
  - ◆ Advice: keep code readable, let the compiler optimize

# Barriers and divergence

- Remember barriers cannot be called in divergent blocks
  - ◆ All threads or none need to call `__syncthreads()`

- If: need to split


```
if(p) {  
    ...  
    // Sync here?  
    ...  
}
```



```
if(p) {  
    ...  
}  
__syncthreads();  
if(p) {  
    ...  
}
```

- Loops: what if trip count depends on data?

```
while(p) {  
    ...  
    // Sync here?  
    ...  
}
```



?

# Barriers instructions

- Barriers with boolean reduction
  - ◆ `__syncthreads_or(p) / __syncthreads_and(p)`  
Synchronize, then returns the boolean OR / AND of all thread predicates p
  - ◆ `__syncthreads_count(p)`  
Synchronize, then returns the number of non-zero predicates
- Loops: what if trip count depends on data?
  - ◆ Loop while at least one thread is still in the loop

```
while(p) {  
    ...  
    // Sync here?  
    ...  
}
```



```
while(__syncthreads_or(p)) {  
    if(p) {  
        ...  
    }  
    __syncthreads();  
    if(p) {  
        ...  
    }  
}
```

# Recap: instruction optimization

---

- Beware of local arrays  
use static indices and loop unrolling
- Use existing math functions
- Keep in mind branch divergence when writing algorithm
  - ◆ But do not end up managing divergence yourself



# Code optimization tools

---

- IDE + profiler: nView
- Profile application: NVIDIA Visual Profiler
- Examine assembly code: cuobjdump
  - ◆ `cuobjdump - -dump-sass`

# OK, I lied

---

- You were told
  - ◆ CPU and GPU have distinct memory spaces
  - ◆ Blocks cannot communicate
  - ◆ You need to synchronize threads inside a block
- This is the least-common denominator across all CUDA GPUs
- This is not true (any more). We now have:
  - ◆ Device-mapped, Unified virtual address space, Unified memory
  - ◆ Global and shared memory atomics
  - ◆ Warp-synchronous programming, warp vote, shuffle
  - ◆ Dynamic parallelism
- Many features are still specific to CUDA and Nvidia GPUs
  - ◆ Should become available in OpenCL eventually...


# Outline

---

- Instruction-level optimization
  - ◆ A few device code features
  - ◆ Warp divergence
- Advanced features
  - ◆ Device-mapped memory, Unified virtual address space, Unified memory
  - ◆ Global and shared memory atomics
  - ◆ Warp-synchronous programming, warp vote, shuffle

# Device-mapped memory

- The GPU can access host memory directly, if explicitly mapped in the device address space
- **Lower bandwidth** but **lower latency** than DMA transfers
  - ◆ Useful for light or sparse memory accesses
- Advantageous for integrated GPUs
- Still different **address spaces**  
Device address is different than host address



# Example

```
// Allocate device-mappable memory
float* a;
cudaHostAlloc((void**)&a, bytes, cudaHostAllocMapped);

// Get device-side pointer
float* da;
cudaHostGetDevicePointer((void**)&d_a, a, 0);

// Write using host pointer on CPU
a[0] = 42;

// Read/write using device pointer on GPU
mykernel<<<grid, block>>>(d_a);

// Wait for kernel completion
cudaDeviceSynchronize();

// Read using host pointer
printf("a=%f\n", a[0]);

// Free memory
cudaFreeHost(a)
```

# Unified virtual address space (UVA)

- Coordinate allocation in CPU and GPU address space
  - ◆ Can tell if an address is on the device, host or both
- `cudaHostAlloc` returns an address valid in both host and device space
- No need to specify direction in `cudaMemcpy`
- Caveat: not a true unified address space
  - ◆ Having an address does not mean you can access the data it points to
  - ◆ Memory allocated with `malloc` still cannot be accessed by GPU
- Requirements: CUDA  $\geq$  4.0, 64-bit OS, CC  $\geq$  2.0
  - ◆ Not available on Windows Vista/7 under WDDM

# Unified memory: the real thing

---

- Allocate memory using `cudaMallocManaged`
  - ◆ The CUDA runtime will take care of the transfers
  - ◆ No need to call `cudaMemcpy` any more
  - ◆ Behaves as if you had a single memory space
- Suboptimal performance: software-managed coherency
  - ◆ Still call `cudaMemcpy` when you can
- Requires CUDA  $\geq 6.0$ , CC  $\geq 3.0$  (preferably 5.x), 64-bit Linux or Windows

# VectorAdd using unified memory

```
int main()
{
    int numElements = 50000;
    size_t size = numElements * sizeof(float);

    float *A, *B, *C;
    cudaMallocManaged((void **)&A, size);
    cudaMallocManaged((void **)&B, size);
    cudaMallocManaged((void **)&C, size);
    Initialize(A, B);

    int blocks = numElements;
    vectorAdd2<<<blocks, 1>>>(A, B, C);
    cudaDeviceSynchronize();
    Display(C);

    cudaFree(A);
    cudaFree(B);
    cudaFree(C);
}
```

- Can I replace `cudaMallocManaged` by `cudaHostAlloc` (with UVA?)
  - ◆ What is the difference?



# Multi-GPU programming

- What if I want to use multiple GPUs on the same machine?
- `cudaGetDeviceCount`, `cudaGetDeviceProperties`  
enumerate devices
- `cudaSetDevice(i)` selects the current device
  - ◆ All following CUDA calls in the thread will concern this device
- Streams are associated with devices
  - ◆ But `cudaStreamWaitEvent` can synchronize with events on other GPU:  
allow inter-GPU synchronization
- Host memory accessible from multiple devices:  
`cudaHostAlloc(..., cudaHostAllocPortable)`
  - ◆ Then call `cudaHostGetDevicePointer` for each GPU

# Peer-to-peer memory copy/access

- Transfer memory between GPUs

- ◆ `cudaMemcpyDeviceToDevice` will not work. Why?


- ◆ Without UVA  
`cudaMemcpyPeer()`  
`cudaMemcpyPeerAsync()`

- ◆ With UVA:  
just plain `cudaMemcpy`

- Direct access to other GPU's memory

- ◆ Check result of `cudaDeviceCanAccessPeer`

- ◆ Call `cudaDeviceEnablePeerAccess` from accessing GPU



# Recap

---

- We have too many mallocs...
  - ◆ `cudaMalloc`
  - ◆ `cudaMallocHost`
  - ◆ `cudaHostAlloc(..., cudaHostAlloc{Portable|Mapped})`
  - ◆ `cudaMallocManaged`
- And too many memcpyys
  - ◆ `cudaMemcpy(dest, src, cudaMemcpy{HostToDevice, DeviceToHost, DeviceToDevice, Default})`
  - ◆ `cudaMemcpyAsync`
  - ◆ `cudaMemcpyPeer[Async]`
- Quizz: do you remember what they do and when we should use them?

# Recap: evolution of the memory model

---

New features: going away from the split memory model

- Device-mapped: GPU can map and access CPU memory
  - ◆ Lower bandwidth than DMA transfers
  - ◆ Higher latency than GPU memory access
- Uniform virtual addressing (CC 2.0):  
synchronize memory space between CPU and GPU
  - ◆ Addresses are unique across the system
  - ◆ No need to specify direction in `cudaMemcpy`
- Unified memory (CC 3.x):  
both CPU and GPU share a managed memory space
  - ◆ Driver manages transfers automatically
  - ◆ Only for memory allocated as managed
  - ◆ Unmanaged + `cudaMemcpy` still useful for optimized transfers

# Outline

---

- Instruction-level optimization
  - ◆ A few device code features
  - ◆ Warp divergence
- Advanced features
  - ◆ Device-mapped memory, Unified virtual address space, Unified memory
  - ◆ Global and shared memory atomics
  - ◆ Warp-synchronous programming, warp vote, shuffle

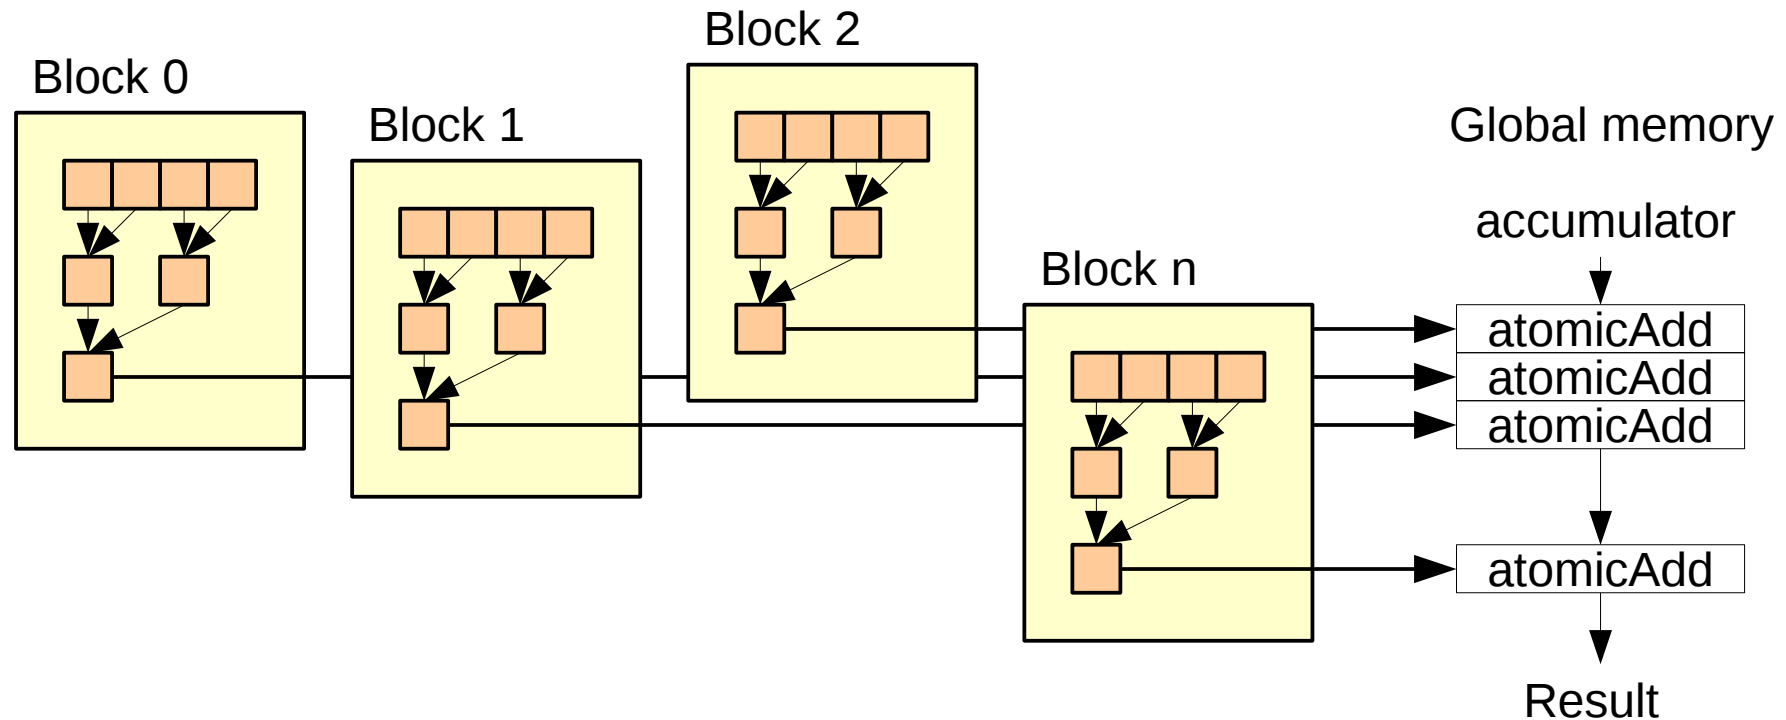
# Atomics

---

- Read, modify, write in one operation
  - ◆ Cannot be mixed with accesses from other thread
- Available operators
  - ◆ Arithmetic: `atomic{Add,Sub,Inc,Dec}`
  - ◆ Min-max: `atomic{Min,Max}`
  - ◆ Synchronization primitives: `atomic{Exch,CAS}`
  - ◆ Bitwise: `atomic{And,Or,Xor}`
- On global memory
  - ◆ From CC 1.1
- On shared memory
  - ◆ From CC 1.2
- Performance impact in case of contention

# Example: reduction

- After local reduction inside each block, use atomics to accumulate the result in global memory



- Complexity?
- Time including kernel launch overhead?

# Floating-point atomics


---

- `atomicAdd` supports single-precision floating-point (float) operands
- Remember floating-point addition is not associative
  - ◆ You will get a different answer depending on the scheduling



# Memory consistency model


- x86 CPUs implement a strong consistency model
  - ◆ Pretend there is a global ordering between memory accesses
- Nvidia GPUs implement a relaxed consistency model
  - ◆ Threads may not see the writes/atomics in the same order



- Need to enforce explicit ordering

# Memory consistency model

- `__threadfence_block`  
`__threadfence`  
`__threadfence_system`  
make all previous writes of the thread visible  
at the block / device / system level




# Outline

---

- Instruction-level optimization
  - ◆ A few device code features
  - ◆ Warp divergence
- Advanced features
  - ◆ Device-mapped memory, Unified virtual address space, Unified memory
  - ◆ Global and shared memory atomics
  - ◆ Warp-synchronous programming, warp vote, shuffle

# Warp-synchronous programming

- We know threads in a warp run synchronously
  - ◆ No need to synchronize them explicitly
- Can use SIMD (PRAM-style) algorithms inside warps
- Example: last steps of a reduction




# Warp-synchronous programming: tools

- We need warp size, and thread ID inside a warp: lane ID
- Official support
  - ◆ Predefined variable: `warpSize`
  - ◆ Lane ID exists in PTX, not in C for CUDA  
Needs to be computed:  
`unsigned int laneId = threadIdx.x % warpSize;`
- Note: as of CUDA 5.0, this is essentially useless
  - ◆ `warpSize` is a variable in PTX, only becomes a constant in SASS
  - ◆ PTX optimizer does not know warp size is a power of 2:  
does not turn `%` into shift
- Often use a `WARP_SIZE` constant hardcoded to 32...

# Warp vote instructions

- `p2 = __all(p1)`  
horizontal AND between the predicates `p1` of all threads in the warp
- `p2 = __any(p1)`  
OR between all `p1`
- `n = __ballot(p)`  
Set bit  $i$  of integer `n` to value of `p` for thread  $i$   
i.e. get bit mask as an integer

Like `__syncthreads_{and,or}` for a warp  
Use: take control decisions for the whole warp  
For CC  $\geq 2.0$



# Manual divergence management

- How to write an if-then-else in warp-synchronous style?  
i.e. without breaking synchronization

- ◆ Using predication:  
execute both sides always

```
if(p) { A(); B(); }  
else { C(); }
```

```
if(p) { A(); }  
// Threads synchronized  
if(p) { B(); }  
// Threads synchronized  
if(!p) { C(); }
```

- ◆ Using vote instructions:  
only execute taken paths  
Skip block if no thread takes it

```
if(__any(p)) {  
    if(p) { A(); }  
    // Threads synchronized  
    if(p) { B(); }  
}  
if(__any(!p)) {  
    // Threads synchronized  
    if(!p) { C(); }  
}
```


- How to write a while loop?

# Shuffle

## Exchange data between lanes

- `__shfl(v, i)`  
Get value of thread  $i$  in the warp
  - ◆ Use: 32 concurrent lookups in a 32-entry table
  - ◆ Use: arbitrary permutation...
- `__shfl_up(v, i) = __shfl(v, tid-i),`  
`__shfl_down(v, i) = __shfl(v, tid+i)`  
Same, indexing relative to current lane
  - ◆ Use: neighbor communication, shift
- `__shfl_xor(v, i) = __shfl(v, tid ^ i)`
  - ◆ Use: exchange data pairwise: “butterfly”


For CC  $\geq 3.0$





# Example: reduction + broadcast

- Naive algorithm



$$a[2*i] \leftarrow a[2*i] + a[2*i+1]$$

$$a[4*i] \leftarrow a[4*i] + a[4*i+2]$$


$$a[8*i] \leftarrow a[8*i] + a[8*i+4]$$

$$a[i] \leftarrow a[0]$$

- Let's rewrite it using shuffle

# Example: reduction + broadcast

- With shuffle



```
ai += __shfl_xor(ai, 1);
```


```
ai += __shfl_xor(ai, 2);
```

```
ai += __shfl_xor(ai, 4);
```

- Exercise: implement complete reduction without `__syncthreads`
  - Hint: use shared memory atomics

# Other example: parallel prefix

- Remember our PRAM algorithm



$\Sigma_{i-j}$  is the sum  $\sum_{k=i-j}^i a_k$

```
s[i] ← a[i]
if i ≥ 1 then
    s[i] ← s[i-1] + s[i]
```


```
if i ≥ 2 then
    s[i] ← s[i-2] + s[i]
```

```
if i ≥ 4 then
    s[i] ← s[i-4] + s[i]
```

```
Step d: if i ≥ 2d then
        s[i] ← s[i-2d] + s[i]
```

# Other example: parallel prefix

- Using warp-synchronous programming



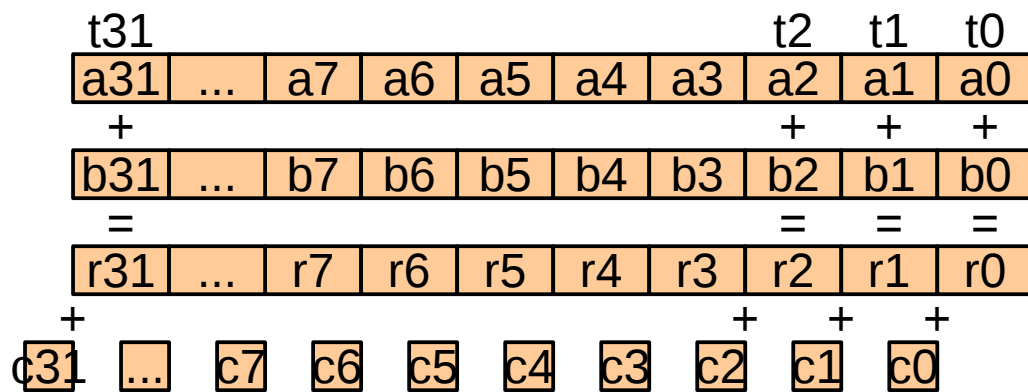
$\Sigma_{i-j}$  is the sum  $\sum_{k=i}^j a_k$

```
s = a;
n = __shfl_up(s, 1);
if(laneid >= 1)
    s += n;
n = __shfl_up(s, 2);
if(laneid >= 2)
    s += n;
n = __shfl_up(s, 4);
if(laneid >= 4)
    s += n;
```

```
for(d = 1; d <= 5; d *= 2) {
    n = __shfl_up(s, d);
    if(laneid >= d)
        s += n;
}
```

# Example: multi-precision addition

- Do an addition on 1024-bit numbers
- Represent numbers as vectors of 32×32-bit
  - ◆ A warp works on a vector
- First step: add elements of the vectors in parallel and recover carries



```
uint32_t a = A[tid],  
          b = B[tid], r, c;  
  
r = a + b;    // Sum  
  
c = r < a;    // Get carry
```


# Second step: propagate carries

- This is a parallel prefix operation

- ◆ We can do it in  $\log(n)$  steps

- But in most cases, one step will be enough

- ◆ Loop until all carries are propagated



```
uint32_t a = A[tid],  
         b = B[tid], r, c;
```

```
r = a + b;    // Sum  
c = r < a;    // Get carry  
while(__any(c)) { // Carry left?  
    c = __shfl_up(c, 1); // Move left  
    if(laneid == 0) c = 0;  
    r = r + c;    // Sum carry  
    c = r < c;    // New carry?  
}  
R[tid] = r;
```

# Takeaway

---

- Two ways to program an SIMT GPU
  - ◆ With independent threads, grouped in warps
  - ◆ With warps operating on vectors
- 3 levels
  - ◆ Blocks in grid: independent tasks, no synchronization
  - ◆ Warps in block: concurrent “threads”, explicitly synchronizable
  - ◆ Threads in warp: implicitly synchronized

# Things we have not talked about

---

- Constant memory
  - ◆ Memory-space that is read-only on the device
  - ◆ Being replaced by cached, read-only access to global memory on recent GPUs
- Texture memory
  - ◆ Cached, read-only memory space optimized for 2D locality
  - ◆ Can unpack compact integer and floating-point encoded data
  - ◆ Can perform filtering: interpolation between data points
- Dynamic parallelism
  - ◆ Starting from CC 3.5, kernels can launch kernels



# Conclusion: trends

---

- GPU: rapidly evolving hardware and software
- Going towards CPU-GPU tighter coupling
  - ◆ On-chip integration
  - ◆ Shared physical memory
  - ◆ Shared virtual memory space
- Most development is going into mobile
  - ◆ Nvidia: Kepler GPUs in Tegra 4 support CUDA
  - ◆ Many GPUs supporting OpenCL:  
ARM Mali, Qualcomm Adreno, Imagination Technologies PowerVR Rogue...
- Still much work to do at the operating system level
  - ◆ GPU currently a second-class citizen
  - ◆ Need to move from CPU+devices model to heterogeneous compute model

# Updated schedule

Week	Course: Tuesday 1:00pm room 2014	Lab: Thursday 1:00pm room 2011
06/10	Programming models	Parallel algorithms
13/10	GPU architecture 1	Know your GPU
20/10	GPU programming	Computing $\ln(2)$ the hard way
27/10	GPU optimization 1	
03/11	GPU optimization 2	Game of life
10/11	Advanced features	
17/11	Lecture by Fernando Pereira	
24/11	Exam	Project
01/12	Project presentations	

# Project

---

- For next week: choose a CUDA programming project
- Work by yourself or in pair
- Ideas of subjects
  - ◆ Image processing: filtering, convolution
  - ◆ Cryptography: brute-force password “recovery”
  - ◆ Data compression
  - ◆ Any proposition?

# References and further reading/watching

---

- CUDA C Programming Guide
- Mark Harris. Introduction to CUDA C.  
<http://developer.nvidia.com/cuda-education>
- David Luebke, John Owens. Intro to parallel programming.  
Online course. <https://www.udacity.com/course/cs344>
- Paulius Micikevicius. GPU Performance Analysis and Optimization. GTC 2012.  
<http://on-demand.gputechconf.com/gtc/2012/presentations/S0514-GTC2012-GPU-Performance-Analysis.pdf>

---

# Pitfalls and fallacies

---

# GPUs are 100x faster than CPUs

---

# GPUs are 100x faster than CPUs

- Wrong
  - ◆ The gap in peak performance (compute and memory) is 10x
  - ◆ In practice, 5x to 8x against optimized CPU code
- Right: you can get a 100x speedup when porting an application to GPU
  - ◆ You get 10x because of the higher GPU throughput
  - ◆ ... and 10x more because you optimized the application  
But spending the same effort on optimization for CPU would also result in 10x improvement
- Right: if you use GPU hardware features not available on CPUs
  - ◆ Texture filtering, elementary functions (exp, log, sin, cos)...

# GPUs have thousands of cores

---



# GPUs have thousands of cores

---

- Official definition of CUDA Core: one execution unit
  - ◆ Right: GPUs have thousands of cores
  - ◆ Under this metric, CPUs have hundreds of cores !
- The closest to a CPU core is a GPU SM
  - ◆ Wrong: 15 to 20 SMs “only” on a GPU
  - ◆ Wide SIMD units inside  
But just x2 to 4x wider than on CPU
- Often-overlooked aspect:  
GPUs have 10s of thousands of threads
  - ◆ But: CPU thread ~ GPU warp
  - ◆ Still thousands of warps

# I can avoid divergence with predication

---

# I can avoid divergence with predication

Do not confuse predication with prediction !

- In a superscalar CPU, branches are predicted
  - ◆ Cost of misprediction ~ pipeline depth  
independent from length of mispredicted path
  - ➔ Inefficient on **short** ifs governed by **random** conditions  
Predication **might** help in this case
- In a GPU, instructions are predicated
  - ◆ Cost of divergence ~ divergent section size
  - ➔ Inefficient on **long** ifs/loops governed by divergent conditions  
Predication **does not** help in this case

# Texture memory

---

- Primarily made for graphics
- Optimized for 2D accesses
  - ◆ Takes advantage of locality in space
- Read-only
- Initialization through specific host functions
  - ◆ `cudaCreateTextureObject()`, `cudaBindTexture2D()`...
- Access through specific device functions
  - ◆ `tex1D()`, `tex2D()`, `tex3D()`...

# Constant memory

---

- Array or variable declared as `__const__`
- Initialized in place or from host code, then read-only
- Single-ported memory: good performance when all threads in a warp access the same word (broadcast)
  - ◆ Similar to shared memory with 1 bank
  - ◆ Generally not worth using with divergent accesses